Suffix Tree (Ukkonen's algorithm)

Wing

February 22, 2019

Table of contents

Trie

An ordered tree data structure used to store a dynamic set or map where the keys are usually strings
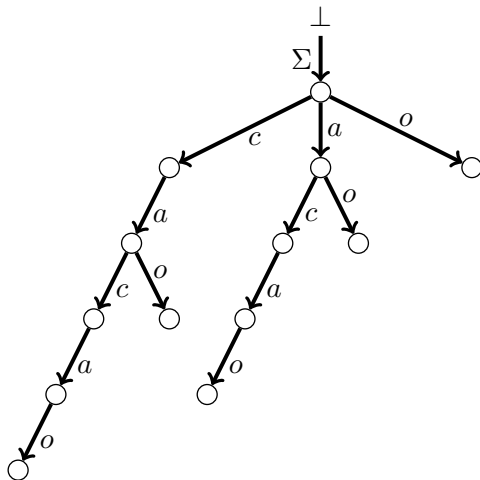
Figure: Suffix Trie for *"cacao"* (Suffix links omitted)

Suffix Trie

- Proposed by Esko Ukkonen (University of Helsinki, Finland)
- An algorithm easier to grasp than the those in the literature at that time
- On-line algorithm: Processes the string symbol by symbol from left to right, and always has the suffix tree for the scanned part of the string ready

String

Let $T = t_1 t_2 \cdots t_n$ be a string over alphabet Σ

Substring

Each string $x : T = uxv$ for some (possibly empty) string u and v is a substring of T

Suffix

$T_i = t_i \cdots t_n$ where $1 \leq i \leq n + 1$

- $T_{n+1} = \epsilon$ is the *empty* suffix

Set of all suffixes of T

$\sigma(T)$

The suffix trie of T is a trie representing $\sigma(T)$

Suffix Trie

Denote suffix trie of T as $STrie(T) = (Q \cup \{\bot\}, root, F, g, f)$

Define such a trie as an augmented deterministic finite-state automaton which has a tree-shaped transition graph representing the trie for $\sigma(T)$

augmented with

- f : suffix function
- $\bot$: auxiliary state

Set Q of the states of $STrie(T)$

The set Q of the states of $STrie(T)$ can be put in a one-to-one correspondence with the substrings of T.

Denote $\bar{x}$ the state that corresponds to a substring x
Shorthand: $\bar{x} \sim x$

- $root \sim \epsilon$
- set F of final states $\sim \sigma(T)$

Transition function g

$$\begin{cases} g(\bar{x}, a) = \bar{y} & \forall \bar{x}, \bar{y} \in Q : y = xa, \text{ where } a \in \Sigma \\ g(\perp, a) = root & \forall a \in \Sigma \end{cases}$$

Suffix function f

$\forall \bar{x} \in Q,$

$$\begin{cases} f(\bar{x}) = \bar{y} & \text{if } \bar{x} \neq root, \text{then } x = ay, a \in \Sigma \\ f(root) = \perp \\ f(\perp) \text{ is undefined} \end{cases}$$

$\perp \sim a^{-1} \; \forall a \in \Sigma$

$a^{-1}a = \epsilon$

Suffix Link

$f(r)$ is the suffix link of state r

Prefix

$T^i = t_1 \cdots t_i$ of T for $0 \le i \le n$

Key observation

How is $STrie(T^i)$ obtained from $STrie(T^{i-1})$?

The suffixes of T^i can be obtained by catenating t_i to the end of each suffix of T^{i-1} and by adding an empty suffix, i.e.

$$\sigma(T^i) = \sigma(T^{i-1})t_i \cup \{\epsilon\}$$

$STrie(T^{i-1})$ accepts $\sigma(T^{i-1})$, to make it accept $\sigma(T^i)$, examine F_{i-1} of $STrie(T^{i-1})$

- $r \in F_{i-1}$ doesn't have a t_i-transition $\Rightarrow$ add transition $r \rightarrow$ new state
- $r \in F_{i-1}$ has a t_i-transition $\Rightarrow$ follow the transition to the old state
- All such states plus $root$ will be F_i of $STrie(T^i)$

How to find states $r \in F_{i-1}$ that get new transitions?

From definition of the suffix function f,
$r \in F_{i-1} \Leftrightarrow r = f^j(\overline{t_1 \cdots t_{i-1}})$ for some $0 \leq j \leq i - 1$

Boundary path

Boundary path of $STrie(T^{i-1})$:
Path starting from deepest state $\overline{t_1 \cdots t_{i-1}}$ of $STrie(T^{i-1})$,
following the suffix links and ending at $\bot$

$\therefore$ All states in F_{i-1} are on the boundary path of $STrie(T^{i-1})$

The boundary path is traversed.

If a state $\bar{z}$ on the boundary path does not have a transition on t_i yet, add a new state $\overline{zt_i}$ and a new transition $g(\bar{z}, t_i) = \overline{zt_i}$

To update f, new states $\overline{zt_i}$ are linked together with new suffix links starting from $\overline{t_1 \cdots t_i}$.
Obviously, this is the boundary path of $STrie(T^i)$

Observation

The traversal over F_{i-1} along the boundary path can be stopped immediately when the first state $\bar{z}$ is found s.t. state $\overline{zt_i}$ (and hence also transition $g(\bar{z}, t_i) = \overline{zt_i}$) already exists.

Let namely $\overline{zt_i}$ already be a state.

Then $STrie(T^{i-1})$ has to contain state $\overline{z't_i}$ and transition $g(z', t_i) = \overline{z't_i} \; \forall z' = f^j(\bar{z}), j \leq 1$.
In other words, if $\overline{zt_i}$ is a substring of T_{i-1} then every suffix of $\overline{zt_i}$ is a substring of T_{i-1}.

Such $\bar{z}$ must exist as $\perp$ is the last state on the boundary path that has the t_i-transition $\forall t_i$

top denotes the state $\overline{t_1 \cdots t_{i-1}}$

Algorithm 1:

---

1 $r \leftarrow top$;
2 **while** $g(r, t_i)$ *is undefined* **do**
3 | create new state r' and new transition $g(r, t_i) = r'$;
4 | **if** $r \neq top$ **then** create new suffix link $f(oldr') = r'$;
5 | $oldr' \leftarrow r'$;
6 | $r \leftarrow f(r)$;
7 create new suffix link $f(oldr') = g(r, t_i)$;
8 $top \leftarrow g(top, t_i)$.

Running Algorithm 1 for $t_i = t_1, t_2, \cdots, t_n$ visits each $\bar{x} \in Q$ once.

Theorem 1

Suffix trie $STrie(T)$ can be constructed in time proportional to the size of $STrie(T)$ which, in the worst case, is $\mathcal{O}(|T|^2)$.

Look at the figure.

Suffix tree

Suffix tree $STree(T)$ of T is a data structure that represents $STrie(T)$ in space linear in the length $|T|$ of T

Construction of Suffix Tree

Explicit states

$Q' \cup \{\bot\}$ is the explicit states of $STrie(T)$

$Q' \subseteq Q$ consists of all branching states and all leaves of $STrie(T)$
By definition, $root$ is included into the branching states

Implicit states

Other states of $STrie(T)$ is the implicit states

Generalized transition

$g'(s, w) = r$ in $STree(T)$ represents the string w spelled out by the transition path in $STrie(T)$ between two explicit states s and r

To save space, the string w is represented as a pair (k, p) of pointers: $t_k \cdots t_p = w$
Then $g'(s, (k, p)) = r$

Such pointers exist because there must be a suffix T_i s.t. the transition path for T_i in $STrie(T)$ goes through s and r

Select the smallest such i, and let k and p point to the substring of this T_i that is spelled out by the transition path from s to r

> **a-transition**
>
> A transition $g'(s, (k, p)) = r$ is called an <u>a-transition</u> if $t_k = a$.

Each s can have at most one a-transition $\forall a \in \Sigma$.

Let $\Sigma = \{a_1, a_2, \ldots, a_m\}$.
$g(\perp, a_j) = root$ is represented as $g(\perp, (-j, -j)) = root$ for $j = 1, \ldots, m$.

Hence suffix tree $STree(T)$ has two components: The tree itself and the string T.

$STree(T)$ is of linear size in $|T|$.

$\because Q'$ has at most $|T|$ leaves (≤ 1 leaf for each nonempty suffix)

$\Rightarrow Q'$ has to contain at most $|T| - 1$ branching states (when $|T| > 1$).

$\therefore$ There can be at most $2|T| - 2$ transitions between the states in Q', each taking a constant space because of using pointers instead of an explicit string.

$\Rightarrow$ In implementation, g' can take $\mathcal{O}(|T|)$ space

Suffix function f'

Let $B \subset Q$ be the set of branching states in $STrie(T)$

$\forall \bar{x} \in B,$
$$\begin{cases} f'(\bar{x}) = \bar{y} & \text{if } \bar{x} \neq root, \text{then } x = ay, a \in \Sigma, \bar{y} \in B \\ f'(root) = \bot \end{cases}$$

f' is well-defined $\because \bar{x} \in B \Rightarrow f'(\bar{x}) \in B$.
These suffix links are explicitly represented. Implicit suffix links are helpful but they are imaginary.

Suffix Tree

Denote suffix tree of T as $STree(T) = (Q' \cup \{\bot\}, root, g', f')$

Reference pair

$r = (s, w)$

Refer to an explicit or implicit state r of a suffix tree by a reference pair (s, w) where

s is some explicit state that is an ancestor of r and

w is the string spelled out by the transitions from s to r in the corresponding suffix trie

Canonical reference pair

A reference pair is canonical if s is the closest ancestor of r (and hence, w is shortest possible)

If r is explicit, canonical reference pair of $r = (r, \epsilon)$

Again, represent string w as a pair (k, p) of pointers s.t.
$t_k \cdots t_p = w$.
Then, reference pair (s, w) gets form $(s, (k, p))$.
$(s, \epsilon) = (s, (p + 1, p))$

Caution

No contraints on k and p as long as w spells the string

It is technically convenient to omit the final states in the definition of a suffix tree.

When final states are necessary, either

- add a symbol ♯ which doesn't occurs in T at the end of T or
- traverse from leaf $\bar{T}$ to $root$ and make all the states on the path explicit

In many applications of $STree(T)$, the start location of each suffix is stored with the corresponding state. Such an augmented tree can be used as an index for finding any substring of T.

The algorithm for constructing $STree(T)$ will be patterned after Algorithm 1.

Now, we make precise what Algorithm 1 does.

Let $s_1 = \overline{t_1 \cdots t_{i-1}}, s_2 = \overline{t_2 \cdots t_{i-1}}, s_3, \ldots, s_i = root, s_{i+1} = \bot$ be the states of $STrie(T^{t-1})$ on the boundary path.

Let j be the smallest index s.t. s_j is not a leaf, and let j' be the smallest index s.t. $s_{j'}$ has a t_i-transition.

As s_1 is a leaf and $\bot$ is a non-leaf that has a t_i-transition, both j and j' are well-defined and $j \leq j'$

Lemma 1

Algorithm 1 adds to $STrie(T^{i-1})$ a t_i-transition $\forall s_h, 1 \le h < j'$, s.t.

- for $1 \le h < j$,
  the new transition expands an old branch of the trie that ends at leaf s_h,
- and for $j \le h < j'$, the new transition initiates a new branch from s_h.

Algorithm 1 does not create any other transitions

Active point

s_j is the <u>active point</u> of $STrie(T^{i-1})$

End point

$s_{j'}$ is the <u>end point</u> of $STrie(T^{i-1})$

These states are present, explicitly or implicitly, in $STree(T^{i-1})$

Lemma 1 says

① Leaf states on the boundary path before the active point s_j get a transition that expands an existing branch of the trie.

② Non-leaf states from the active point s_j to end point $s_{j'}$ ($s_{j'}$ excluded) get a transition that initiates a new branch

Interpret in terms of suffix tree $STree(T^{i-1})$.

Transitions from ① that expand an existing branch is implemented by updating the right pointer of each transition that represents the branch.

Let $g'(s, (k, i - 1)) = r$ be such a transition.
The right pointer has to point to the last position $i - 1$ of T^{i-1}.
$\because r$ is a leaf $\Rightarrow$ a path leading to r has to spell out a suffix of T^{i-1} that does not occur elsewhere in T^{i-1}.
$\therefore$ Updated transition is $g'(s, (k, i)) = r$.

This only makes the string spelled out by the transition longer but does not change the states s and r. Making all such updates would take too much time. We use a trick for this.

Open transition

Any transition of $STree(T^{i-1})$ leading to a leaf

Open transitions are represented as $g'(s, (k, \infty)) = r$
Symbols ∞ can be replaced by $n = |T|$ after completing $STree(T)$

This way, transitions from $\textcircled{1}$ is automatically done.

For transitions from ②,
we need to find all s_h, $j \leq h < j'$, but s_h might not be explicit.

Let $h = j$ and let (s, w) be the <u>canonical reference pair</u> for s_h (the active point).

s_h is on the boundary path of $STrie(T^{i-1})$

$\Rightarrow w$ is a suffix of T^{i-1}

$\Rightarrow (s, w) = (s, (k, i-1))$ for some $k \leq i$

We need to create a new branch starting from the state $(s, (k, i - 1))$.

First, if $(s, (k, i - 1))$ is the end point, then done.
Otherwise, $s_h = (s, (k, i - 1))$ has to be explicit in order to create a new branch from there.

If s_h is not explicit, create the explicit state s_h by splitting the transition that contains the corresponding implicit state.
After that, a t_i-transition from s_h is created which is $g'(s_h, (i, \infty)) = s_{h'}$ where $s_{h'}$ is a new leaf.
Moreover, suffix link $f'(s_h)$ is added if s_h is created by splitting a transition.

Next the construction proceeds to s_{h+1}.

Reference pair for $s_h = (s, (k, i-1))$
$\Rightarrow$ canonical reference pair for $s_{h+1} = canonize(f'(s), (k, i-1))$
where $canonize$ makes the pair canonical by updating the state and the left pointer.

Repeat until $s_{j'}$ is found.

update returns a reference pair for the end point $s_{j'}$ (only the state and the left pointer as right pointer is always $i - 1$)

Procedure update$(s, (k, i))$

$(s, (k, i - 1))$ the <u>canonical reference pair</u> for the active point;

1 $oldr \leftarrow root$;

2 $(end\text{-}point, r) \leftarrow test\text{-}and\text{-}split(s, (k, i - 1)), t_i)$;

3 **while not** $(end\text{-}point)$ **do**

4   create new transition $g'(r, (i, \infty)) = r'$ where r' is a new state;

5   **if** $oldr \neq root$ **then** create new suffix link $f'(oldr) = r'$;

6      $oldr \leftarrow r$;

7      $(s, k) \leftarrow canonize(f'(s), (k, i - 1))$;

8      $(end\text{-}point, r) \leftarrow test\text{-}and\text{-}split(s, (k, i - 1), t_i)$;

9 **if** $oldr \neq root$ **then** create new suffix link $f'(oldr) = s'$;

10 **return** (s, k).

---

test-and-split returns ('is end point?', explicit state)

Procedure test-and-split(s, $(k,\ p)$, t)

$(s, (k, p))$ is canonical

1   **if** $k \leq p$ **then**

2   Let $g'(s, (k', p')) = s'$ be the t_k-transition from s;

3   **if** $t = t_{k'+p-k+1}$ **then** **return** (**true**, s)

4   **else**

5   replace the t_k-transition above by transitions
        $g'(s, (k', k' + p - k)) = r$ and
        $g'(r, (k' + p - k + 1, p')) = s'$ where r is a new state;

6   **return** (**false**, r)

7   **else**

8   **if** *there is no t-transition from s* **then** **return** (**false**, s)

9   **else** **return** (**true**, s).

Procedure $test\text{-}and\text{-}split$ benefits from that $(s, (k, p))$ is canonical: The answer to the end point test can be found in constant time by considering only one transition from s

---

Procedure canonize(s, $(k,\ p)$)

---

1   **if** $p < k$ **then** **return** (s, k)

2   **else**

3   find the t_k-transition $g'(s, (k', p')) = s'$ from s;

4   **while** $p' - k' \leq p - k$ **do**

5           $k \leftarrow k + p' - k' + 1$;

6           $s \leftarrow s'$;

7   **if** $k \leq p$ **then** find the t_k-transition $g'(s, (k', p')) = s'$
          from s ;

8   **return** (s, k).

At return, k can become k', $k' \geq k$.

To continue the construction for the next text symbol t_{i+1}, the active point of $STree(T^i)$ has to be found

Fact 1

s_j is the active point of $STree(T^{i-1})$
$\Leftrightarrow s_j = \overline{t_j \cdots t_{i-1}}$ where $t_j \cdots t_{i-1}$ is the longest suffix of T^{i-1} that occurs at least twice in T^{i-1}

∵ By the construction process,

- $t_j \cdots t_{i-1}$ has occured before $\Leftrightarrow s_j$ is not a leaf and
- j is smallest meaning $t_j \cdots t_{i-1}$ is longest

Fact 2

$s_{j'}$ is the end point of $STree(T^{i-1})$

$\Leftrightarrow s_{j'} = \overline{t_{j'} \cdots t_{i-1}}$ where $t_{j'} \cdots t_{i-1}$ is the longest suffix of T^{i-1}

s.t. $t_{j'} \cdots t_{i-1} t_i$ is a substring of T^{i-1}

$\because$

- by definition of end point and
- j is smallest meaning $t_j \cdots t_{i-1}$ is longest

Combining the previous 2 facts gives

$s_{j'}$ is the end point of $STree(T^{i-1})$
$\Rightarrow t_{j'} \cdots t_{i-1} t_i$ is the longest suffix of T^i that occurs at least twice in T^i
$\Leftrightarrow$ state $g'(s_{j'}, t_i)$ is the active point of $STree(T^i)$

$\because t_{j'} \cdots t_{i-1} t_i$ is a substring of T^{i-1}
$\Rightarrow t_{j'} \cdots t_{i-1} t_i$ occurs at least twice in T^i

Lemma 2

$(s, (k, i - 1))$ is reference pair of the end point $s_{j'}$ of $STree(T^{i-1})$
$\Rightarrow (s, (k, i))$ is a reference pair of the active point of $STree(T^i)$

Algorithm 2: Construction of $STree(T)$ for string T = $t_1 t_2 \cdots \sharp$ in alphabet $\Sigma = \{t_{-1}, \ldots, t_{-m}\}$;
$\sharp$ is the end marker not appearing elsewhere in T.

---

1   create states $root$ and $\perp$;
2   **for** $j \leftarrow 1, \ldots, m$ **do**   create transition $g'(\perp, (-j, -j)) = root$;
3   create suffix link $f'(root) = \perp$;
4   $s \leftarrow root$; $k \leftarrow 1$; $i \leftarrow 0$;
5   **while** $t_{i+1} \neq \sharp$ **do**
6       $i \leftarrow i + 1$;
7       $(s, k) \leftarrow update(s, (k, i))$;
8       $(s, k) \leftarrow canonize(s, (k, i))$;

Steps 7-8 are based on Lemma 2

Theorem 2

Algorithm 2 constructs the suffix tree $STree(T)$ for a string $T = t_1 \cdots t_n$ on-line in time $\mathcal{O}(n)$

Proof:

The algorithm constructs $STree(T)$ through intermediate trees $STree(T^0)$, $STree(T^1)$,..., $STree(T^n) = STree(T)$.
It is on-line as to construct $STree(T^i)$ it only needs access to the first i symbols of T.

For the running time analysis,
we divide the time requirement into two components, both turn out
to be $\mathcal{O}(n)$.

◇1◇ Total time for procedure *canonize*

◇2◇ The rest: The time for

repeatedly traversing the suffix link path from the present
active point to the end point and

creating the new branches by update and

then finding the next active point by taking a transition from
the end point

Visted States

Call the states (reference pairs) on these paths the <u>visited states</u>

$\langle 2 \rangle$ takes time proportional to the total number of the visited states
$\because$ the operations at each such state
(create an explicit state and a new branch,
follow an explicit or implicit suffix link,
test for the end point)
at each such state can be implemented in constant time as $canonize$ is excluded.

(To be precise, this also requires that $|\Sigma|$ is bounded independently of n.)

Let r_i be the active point of $STree(T^i)$ for $0 \leq i \leq n$.
The visited states between r_{i-1} and r_i are on a path that consists of some suffix links and one t_i-transition

Depth of a state

The depth of a state is the length of the string spelled out on the transition path from $root$

Taking a suffix link decreases the depth of the current state by 1.
$\therefore$ The number of the visited states (including r_{i-1}, excluding r_i) on the path is $depth(r_{i-1}) - depth(r_i) + 2$,

and their total number is $\sum_{i=1}^{n}(depth(r_{i-1}) - depth(r_i) + 2) = depth(r_0) - depth(r_n) + 2n \leq 2n$.
This implies $\langle 2 \rangle$ takes time $\mathcal{O}(n)$

The time spent by each execution of $canonize$ is $\mathcal{O}(a + bq)$
where a and b are constants and
q is the number of executions of the body of the loop in steps 5-7
of $canonize$.
∴ The total time spent by $canonize$ in all calls =
$\mathcal{O}($number of the calls of $canonize$ +
the total number of the executions of the body of the loop$)$

There are $\mathcal{O}(n)$ calls as there is one call for each visited state (either in step 6 of $update$ or directly in step 8 of Alg. 2.).

Each execution of the body deletes a nonempty string from the left end of string $w = t_k \cdots t_p$

String w can grow during the whole process only in step 6^1 of Alg.2 which catenates t_i for $i = 1, \ldots, n$ to the right end of w.

$\therefore$ a non-empty deletion is possible at most n times.

$\Rightarrow \langle 1 \rangle$ takes time $\mathcal{O}(n)$.

---

[1]my correction

Original Paper:
https://www.cs.helsinki.fi/u/ukkonen/
SuffixT1withFigs.pdf

Wikipedia for the definition of Trie